

Implement a Basic Driving Agent

Q: Observe what you see with the agent's behavior as it takes random actions. Does the smartcab eventually make it to the destination? Are there any other interesting observations to note?

A:

The smartcab is not smart enough and does not make it to the destination almost all the time. Sometimes it can arrive at the destination luckily with its random actions.

The smartcab doesn't wait patiently at intersections and often goes away from the destination. That is natural because the smartcab chooses his action randomly from the 4 choices i.e. None, forward, left and right, and the chance of choosing None and doing nothing and waiting is less than that of choosing the others and making moves.

Inform the Driving Agent

Q: What states have you identified that are appropriate for modeling the smartcab and environment? Why do you believe each of these states to be appropriate for this problem?

A:

I identified that the states below are appropriate for modeling the smartcab.

Next waypoint information from the route planner:

My smartcab can learn the direction it should go from this information. Without this information, it would not know the right direction to make it to its destination.

Traffic light:

With the information of the next waypoint, the smartcab can approach the destination. But it must obey the traffic rules. So it also needs the information of traffic light (red or green).

Oncoming traffic, Left traffic:

To avoid accidents with other cabs, it should care oncoming traffics and left traffics. When it wants to make left turn at intersections, it must be sure that there is no oncoming traffic. Similarly, when it wants to make right turn, it must be sure that oncoming traffic is approaching from its left through the intersection. I think it doesn't need to care the oncoming traffic approaching from its right. That's because there is no chance of my smartcab crashing into its right traffic unless it violates the traffic rules and goes through the intersections with the traffic light red.

Implement a Q-Learning Driving Agent

Q: What changes do you notice in the agent's behavior when compared to the basic driving agent when random actions were always taken? Why is this behavior occurring?

A:

Compared to the basic driving agent that always takes random actions, the agent with Q-Learning gradually learns how to make it to the destination. At first several trials, it doesn't know which way to go and takes almost random actions. It can't even approach the destination appropriately, much less arrive there. But as the trial goes on, it learns how to approach the destination and finally becomes to be able to arrive there.

Q-Learning makes this change in the agent's behavior. The agent learns the state it is in currently and the reward it takes from its action, and becomes to be able to select the appropriate action to get the maximum reward.

Improve the Q-Learning Driving Agent

Q: Report the different values for the parameters tuned in your basic implementation of Q-Learning. For which set of parameters does the agent perform best? How well does the final driving agent perform?

A:

I tried different values for the parameters alpha, gamma and epsilon and created the table below. The columns from 1st to 10th represent the number of simulations (, each of which consist of 100 trials) and the each value represents how many time the agent was able to get to the destination. For example, with alpha = 0.1, gamma = 0.9 and epsilon = 0.1, the agent was able to reach the destination 52 times in the 1st simulation. Summing up the values from 1st to 10th, I got the average number of success 69.6 for alpha = 0.1, gamma = 0.9 and epsilon = 0.1.

alpha	gamma	epsilon	avg. success	simulation:	1st	2nd	3rd	4th	5th	6th	7th	8th	9th	10th
0.1	0.9	0.1	69.6		52	80	89	86	25	79	83	85	84	33
0.2	0.1	0.1	85.6		78	89	88	87	81	91	85	92	79	86
0.2	0.2	0.1	85.7		90	74	89	90	87	88	94	66	85	94
0.2	0.3	0.1	81.2		89	93	81	55	85	90	79	70	81	89
0.2	0.4	0.1	82.5		92	79	80	75	86	85	86	82	86	74
0.2	0.5	0.1	78.8		85	84	84	75	83	56	78	78	86	79
0.2	0.8	0.1	63		51	67	76	60	57	45	40	78	77	79
0.2	0.9	0.1	76.1		77	80	74	84	73	72	85	79	54	83
0.3	0.9	0.1	74.2		79	71	85	48	78	80	71	79	76	75
0.4	0.9	0.1	62.4		66	63	65	47	72	64	48	65	64	70
0.5	0.9	0.1	66.2		65	74	58	53	66	73	69	56	70	78
0.5	0.9	0.2	58.7		61	63	56	61	56	62	47	61	63	57
0.5	0.9	0.3	51.5		61	47	44	55	47	39	58	49	55	60
0.5	0.9	0.5	41.5		45	30	42	47	36	40	46	38	48	43
0.6	0.9	0.1	63.6		64	64	59	67	62	76	65	54	63	62
0.9	0.9	0.5	36.7		40	40	37	31	34	42	31	39	25	48

The set of (Alpha = 0.2, gamma = 0.2, epsilon = 0.1) gave the agent the best result (the average number of success is 85.7).

I'm particularly convinced that the low epsilon values brought the best result. With high epsilons, the agent takes random actions too much and cannot exploit what it learned. Not able to learn from its experience, the agent has little chance of reaching the destination.

Q: Does your agent get close to finding an optimal policy, i.e. reach the destination in the minimum possible time, and not incur any penalties? How would you describe an optimal policy for this problem?

A:

I describe the optimal policy as follows:

First, checking the next waypoint. The agent knows which way to go and determines the best action to approach the next waypoint directly and immediately. The agent only thinks about the shortest path to the destination. I think checking the next waypoint is a first priority to reach the destination in the minimum possible time.

Then checking the traffic light. The agent ensures that the best action it determines is valid and it should obey the traffic rules. If the action is against the rules i.e. going through the intersections with red light or making left turn at the intersection, the agent turns down the action and selects the action None and waits until the light turns green.

Finally, ensuring the light is green, the agent cares the other agents not to cause traffic accidents. When it becomes sure that there is no traffic to cause accidents and it can make a move safely at the intersection , it determines the action which enables it to get to the next waypoint immediately.